

WORLD-MODEL UNCERTAINTY BLOAT: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot world-model reliability. The motivating bet is: Expose cases where uncertainty mechanisms hide missing mechanics instead of repairing them.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot world-model reliability, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b;c; 2021a;b; 2026d). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.552	0.448
observed_only_wam	hidden_mode_shift	0.163	0.838
augmentation_as_data_wam	nominal	0.560	0.440
augmentation_as_data_wam	hidden_mode_shift	0.268	0.732
scalar_uncertainty_wam	nominal	0.675	0.325
scalar_uncertainty_wam	hidden_mode_shift	0.320	0.680
proposed_branch_mechanism	nominal	0.775	0.225
proposed_branch_mechanism	hidden_mode_shift	0.542	0.458

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

World-Model Uncertainty Bloat argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Zero-shot uncertainty-aware deployment of simulation trained policies on real-world robots, 2021a. <http://arxiv.org/abs/2112.05299v1>.
- Bayesian controller fusion: Leveraging control priors in deep reinforcement learning for robotics, 2021b. <http://arxiv.org/abs/2107.09822v3>.
- An investigation into uncertainty quantification of shallow foundation failure mechanisms in horizontally stratified layered soil strata, 2026a. <https://doi.org/10.3390/app16063051>.
- Vls: Steering pretrained robot policies via vision-language models, 2026b. <http://arxiv.org/abs/2602.03973v1>.
- Uncertainty-calibrated safety gating for vision-language- action manipulation under domain shift: Reliability gains and intervention-efficiency trade-offs, 2026c. <https://doi.org/10.3390/s26103140>.
- Worldbench: Disambiguating physics for diagnostic evaluation of world models, 2026d. <http://arxiv.org/abs/2601.21282v1>.